

Elijah Lin

11775718

Started off by creating `proc_list`, a struct that will store information about processes such as the PID and the CPU time. Had to create a `cpu time to jiffies` converter, to change nanoseconds to jiffies. Then created read and write functions to read in the information about the processes. After reading in the process's information, it's returned to user space.

I implemented a timer function to schedule periodic updates of the CPU times for the processes. Workqueue function to handle background processing. Then there's a timer callback function for when the timer expires then queues work to be done in the background.

Initialize function creates the `proc` directory as well as sets up the timer, the work queue, and the necessary locks. The initialize function acts as the module's loading entry point and when it's finished the exit function is called to remove the status and the `proc` as well as free memory.

The `userapp.c` has a `register process` function that just does that, registering processes. Main is essentially a loop that runs indefinitely until start time exceeds the expire time. Once it does loop is broken and the program exits.

The next page has two images, the top represents the two-process case and the bottom represents the one process case.

